

# Statlog DATASET

already preprocessed the dataset.

```
In [13]: from pyspark.ml.feature import HashingTF, IDF, Tokenizer
from pyspark.sql import SparkSession
from pyspark.sql.functions import monotonically_increasing_id
from pyspark.sql import Row
import csv
spark = SparkSession.builder.getOrCreate()
f = open("pima/australian.dat")
reader = csv.reader(f, delimiter=' ')
ww = []
for w in reader:
    ww.append(w)

data = map(lambda p: Row(label=int(p[0]),
a_1=float(p[1]),a_2=float(p[2]),
label_2=int(p[3]),label_3=int(p[4]),label_
_4=int(p[5]),
a_3=float(p[6]),label_5 =int(p[7]),label_
6=int(p[8]),
a_4=float(p[9]),label_7=int(p[10]),label_
8=int(p[11]),
a_5=float(p[12]),a_6=float(p[13]),label_9=int(p[14]))
,ww)
data = spark.createDataFrame(data)
f.close()
```

```
In [14]: data.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|  a_1|  a_2|  a_3|  a_4|  a_5|  a_6|label|label_2|label_3|label_4|label_5|label_6|label_7|label_8|label_9|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|22.08|11.46|1.585| 0.0|100.0|1213.0| 1| 2| 4| 4|
| 0| 0| 1| 2| 0|
|22.67| 7.0|0.165| 0.0|160.0| 1.0| 0| 2| 8| 4|
| 0| 0| 0| 2| 0|
|29.58| 1.75| 1.25| 0.0|280.0| 1.0| 0| 1| 4| 4|
| 0| 0| 1| 2| 0|
|21.67| 11.5| 0.0|11.0| 0.0| 1.0| 0| 1| 5| 3|
| 1| 1| 1| 2| 1|
|20.17| 8.17| 1.96|14.0| 60.0| 159.0| 1| 2| 6| 4|
| 1| 1| 0| 2| 1|
|15.83|0.585| 1.5| 2.0|100.0| 1.0| 0| 2| 8| 8|
| 1| 1| 0| 2| 1|
|17.42| 6.5|0.125| 0.0| 60.0| 101.0| 1| 2| 3| 4|
| 0| 0| 0| 2| 0|
|58.67| 4.46| 3.04| 6.0| 43.0| 561.0| 0| 2| 11| 8|
| 1| 1| 0| 2| 1|
|27.83| 1.0| 3.0| 0.0|176.0| 538.0| 1| 1| 2| 8|
| 0| 0| 0| 2| 0|
|55.75| 7.08| 6.75| 3.0|100.0| 51.0| 0| 2| 4| 8|
| 1| 1| 1| 2| 0|
| 33.5| 1.75| 4.5| 4.0|253.0| 858.0| 1| 2| 14| 8|
| 1| 1| 1| 2| 1|
|41.42| 5.0| 5.0| 6.0|470.0| 1.0| 1| 2| 11| 8|
| 1| 1| 1| 2| 1|
|20.67| 1.25|1.375| 3.0|140.0| 211.0| 1| 1| 8| 8|
| 1| 1| 1| 2| 0|
|34.92| 5.0| 7.5| 6.0| 0.0|1001.0| 1| 2| 14| 8|
| 1| 1| 1| 2| 1|
|58.58| 2.71|2.415| 0.0|320.0| 1.0| 1| 2| 8| 4|
| 0| 0| 1| 2| 0|
|48.08| 6.04| 0.04| 0.0| 0.0|2691.0| 1| 2| 4| 4|
| 0| 0| 0| 2| 1|
|29.58| 4.5| 7.5| 2.0|330.0| 1.0| 1| 2| 9| 4|
| 1| 1| 1| 2| 1|
|18.92| 9.0| 0.75| 2.0| 88.0| 592.0| 0| 2| 6| 4|
| 1| 1| 0| 2| 1|
| 20.0| 1.25|0.125| 0.0|140.0| 5.0| 1| 1| 4| 4|
| 0| 0| 0| 2| 0|
|22.42|5.665|2.585| 7.0|129.0|3258.0| 0| 2| 11| 4|
| 1| 1| 0| 2| 1|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
only showing top 20 rows
```

## Vectorize the Data into feature

```
In [15]: from pyspark.ml.feature import VectorAssembler
label = ["label"]
assembler = VectorAssembler(
    inputCols=[x for x in data.columns if x not in label],
    outputCol='features')
data = assembler.transform(data)
data.show()
```

a_1	a_2	a_3	a_4	a_5	a_6	label	label_2	label_3	label_4	label_5	label_6	label_7	label_8	label_9	features
22.08	11.46	1.585	0.0	100.0	1213.0	1	2	4	4						
0	0	1	2	0	[22.08,11.46,1.58...										
22.67	7.0	0.165	0.0	160.0	1.0	0	2	8	4						
0	0	0	2	0	[22.67,7.0,0.165,...										
29.58	1.75	1.25	0.0	280.0	1.0	0	1	4	4						
0	0	1	2	0	[29.58,1.75,1.25,...										
21.67	11.5	0.0	11.0	0.0	1.0	0	1	5	3						
1	1	1	2	1	[21.67,11.5,0.0,1...										
20.17	8.17	1.96	14.0	60.0	159.0	1	2	6	4						
1	1	0	2	1	[20.17,8.17,1.96,...										
15.83	0.585	1.5	2.0	100.0	1.0	0	2	8	8						
1	1	0	2	1	[15.83,0.585,1.5,...										
17.42	6.5	0.125	0.0	60.0	101.0	1	2	3	4						
0	0	0	2	0	[17.42,6.5,0.125,...										
58.67	4.46	3.04	6.0	43.0	561.0	0	2	11	8						
1	1	0	2	1	[58.67,4.46,3.04,...										
27.83	1.0	3.0	0.0	176.0	538.0	1	1	2	8						
0	0	0	2	0	[27.83,1.0,3.0,0....										
55.75	7.08	6.75	3.0	100.0	51.0	0	2	4	8						
1	1	1	2	0	[55.75,7.08,6.75,...										
33.5	1.75	4.5	4.0	253.0	858.0	1	2	14	8						
1	1	1	2	1	[33.5,1.75,4.5,4....										
41.42	5.0	5.0	6.0	470.0	1.0	1	2	11	8						
1	1	1	2	1	[41.42,5.0,5.0,6....										
20.67	1.25	1.375	3.0	140.0	211.0	1	1	8	8						
1	1	1	2	0	[20.67,1.25,1.375...										
34.92	5.0	7.5	6.0	0.0	1001.0	1	2	14	8						
1	1	1	2	1	[34.92,5.0,7.5,6....										
58.58	2.71	2.415	0.0	320.0	1.0	1	2	8	4						
0	0	1	2	0	[58.58,2.71,2.415...										
48.08	6.04	0.04	0.0	0.0	2691.0	1	2	4	4						
0	0	0	2	1	[48.08,6.04,0.04,...										
29.58	4.5	7.5	2.0	330.0	1.0	1	2	9	4						
1	1	1	2	1	[29.58,4.5,7.5,2....										
18.92	9.0	0.75	2.0	88.0	592.0	0	2	6	4						
1	1	0	2	1	[18.92,9.0,0.75,2...										
20.0	1.25	0.125	0.0	140.0	5.0	1	1	4	4						
0	0	0	2	0	[20.0,1.25,0.125,...										
22.42	5.665	2.585	7.0	129.0	3258.0	0	2	11	4						
1	1	0	2	1	[22.42,5.665,2.58...										

only showing top 20 rows

## Split the Data

```
In [17]: splits = data.select("label", "features").randomSplit([0.8, 0.2], 1234)
        train = splits[1]
        test = splits[0]
```

## Use NaiveBayes method to build a model

```
In [18]: from pyspark.ml.classification import NaiveBayes
        from pyspark.ml.evaluation import MulticlassClassificationEvaluator

        nb = NaiveBayes()
        model = nb.fit(train)
        predictions = model.transform(test)

        evaluator = MulticlassClassificationEvaluator(labelCol="label", predictionCol="prediction",
                                                    metricName="accuracy")
        accuracy = evaluator.evaluate(predictions)
        print("Test set accuracy = " + str(accuracy))

Test set accuracy = 0.388791593695
```

## Use DecisionTree method to build a model

```
In [19]: from pyspark.ml.classification import DecisionTreeClassifier
        from pyspark.ml.evaluation import MulticlassClassificationEvaluator

        dt = DecisionTreeClassifier()

        model = dt.fit(train)

        predictions = model.transform(test)

        evaluator = MulticlassClassificationEvaluator(labelCol="label", predictionCol="prediction", metricName="accuracy")
        accuracy = evaluator.evaluate(predictions)
        print("Test set accuracy = " + str(accuracy))

Test set accuracy = 0.647985989492
```

## Use RandomForest method to build a model

```
In [20]: from pyspark.ml.classification import RandomForestClassifier
from pyspark.ml.evaluation import MulticlassClassificationEvaluator

rf = RandomForestClassifier()

model = rf.fit(train)

predictions = model.transform(test)

evaluator = MulticlassClassificationEvaluator(labelCol="label",predictionCol="prediction",metricName="accuracy")
accuracy = evaluator.evaluate(predictions)
print("Test set accuracy of RandomForest= " + str(accuracy))

Test set accuracy of RandomForest= 0.647985989492
```

## Summary

Naive Bayes model fails when the input data is very independent. And the classifiers used in this experiment give out poor performance on continuous data.